



# Introduction to SONIC + Triton (ML inference as a Service)


Yongbin Feng, Kevin Pedro (Fermilab)

Machine Learning HATS@LPC



July 12th, 2021

# Introduction

- Now that you have trained a neural network. **How to deploy the model and run a large scale production (**inference**) on some datasets?**
  - ❖ CMSSW supports TF1, and started supporting TF2 since CMSSW\_11\_1\_X ([cmssw#28711](#), [cmsdist#5525](#)), and ONNX, so that you can run direct inference on CPUs with TF1, TF2, and ONNX models.
  - ❖ What if you have a Pytorch model instead?
    - ➔ Converting it to an ONNX model and run inference with CMSSW is an option
  - ❖ What if you have a GPU (or some other coprocessors), and you want to accelerate the inference with it?
    - ➔ What if the GPU is remote, not directly connected with the CPU clusters?



ONNX  
RUNTIME




# Introduction




- Now that you have trained a neural network. **How to deploy the model and run a large scale production ('inference') on some datasets?**

❖ You can use **SONIC + Triton**, which will help you solve problems such as:

- ❖ What if you have a Pytorch model instead?
  - ➔ Converting it to an ONNX model and run inference with CMSSW is an option
- ❖ What if you have a GPU (or some other coprocessors), and you want to accelerate the inference with it?
  - ➔ What if the GPU is remote, not directly connected with the CPU clusters?



ONNX  
RUNTIME

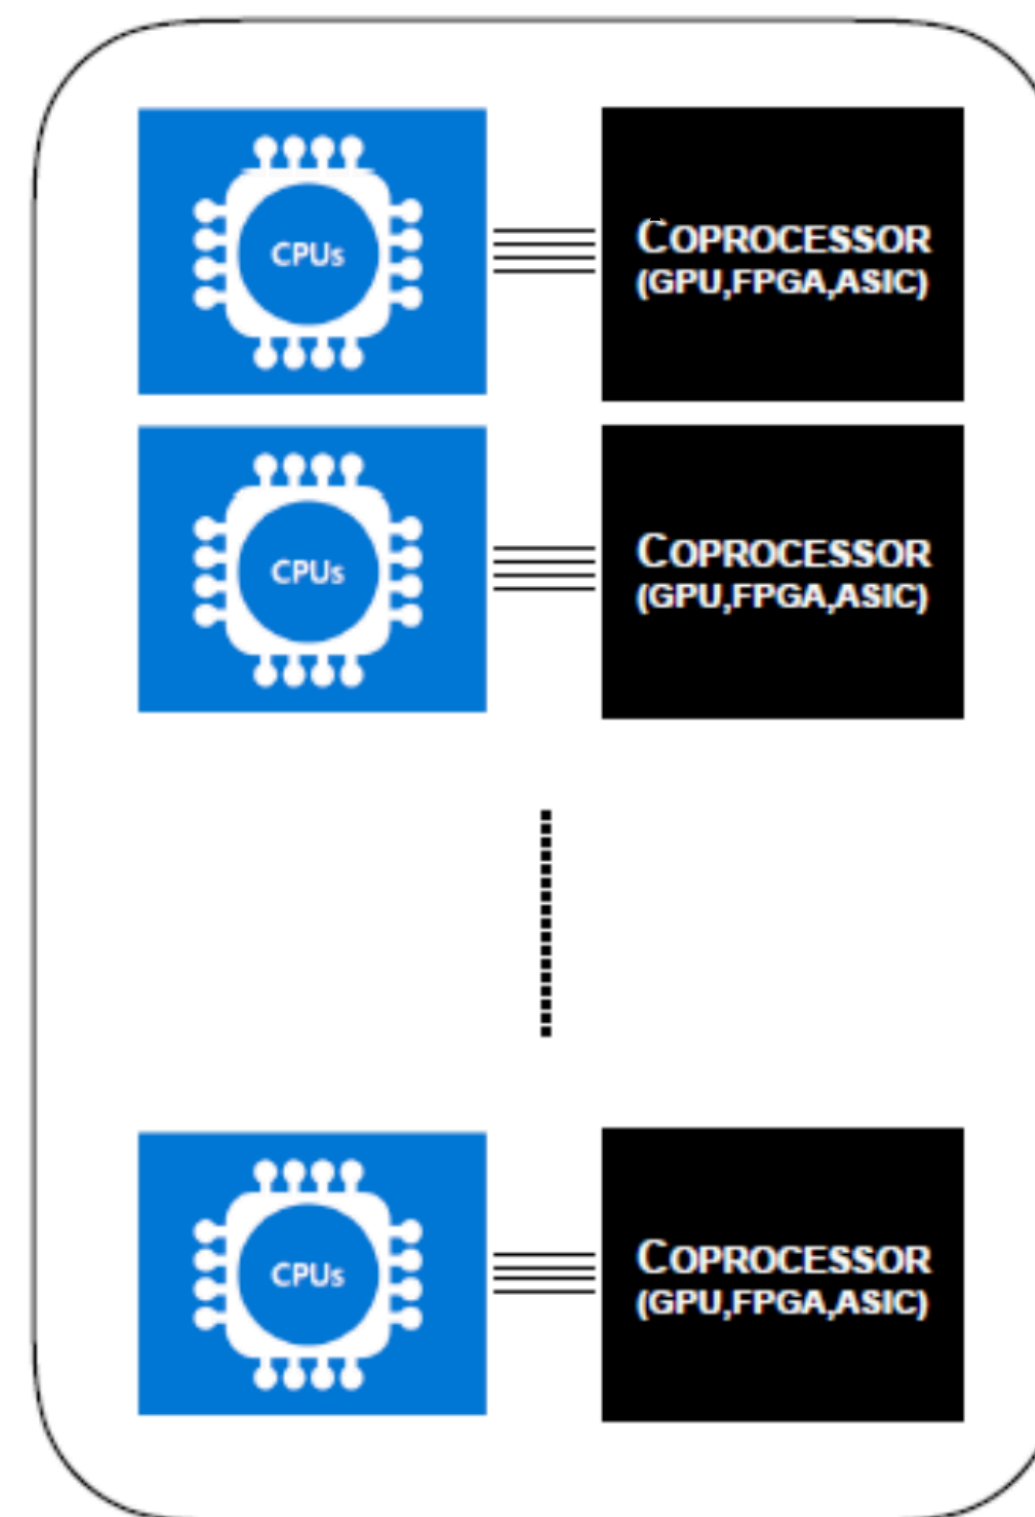


# ML Inference Infrastructure

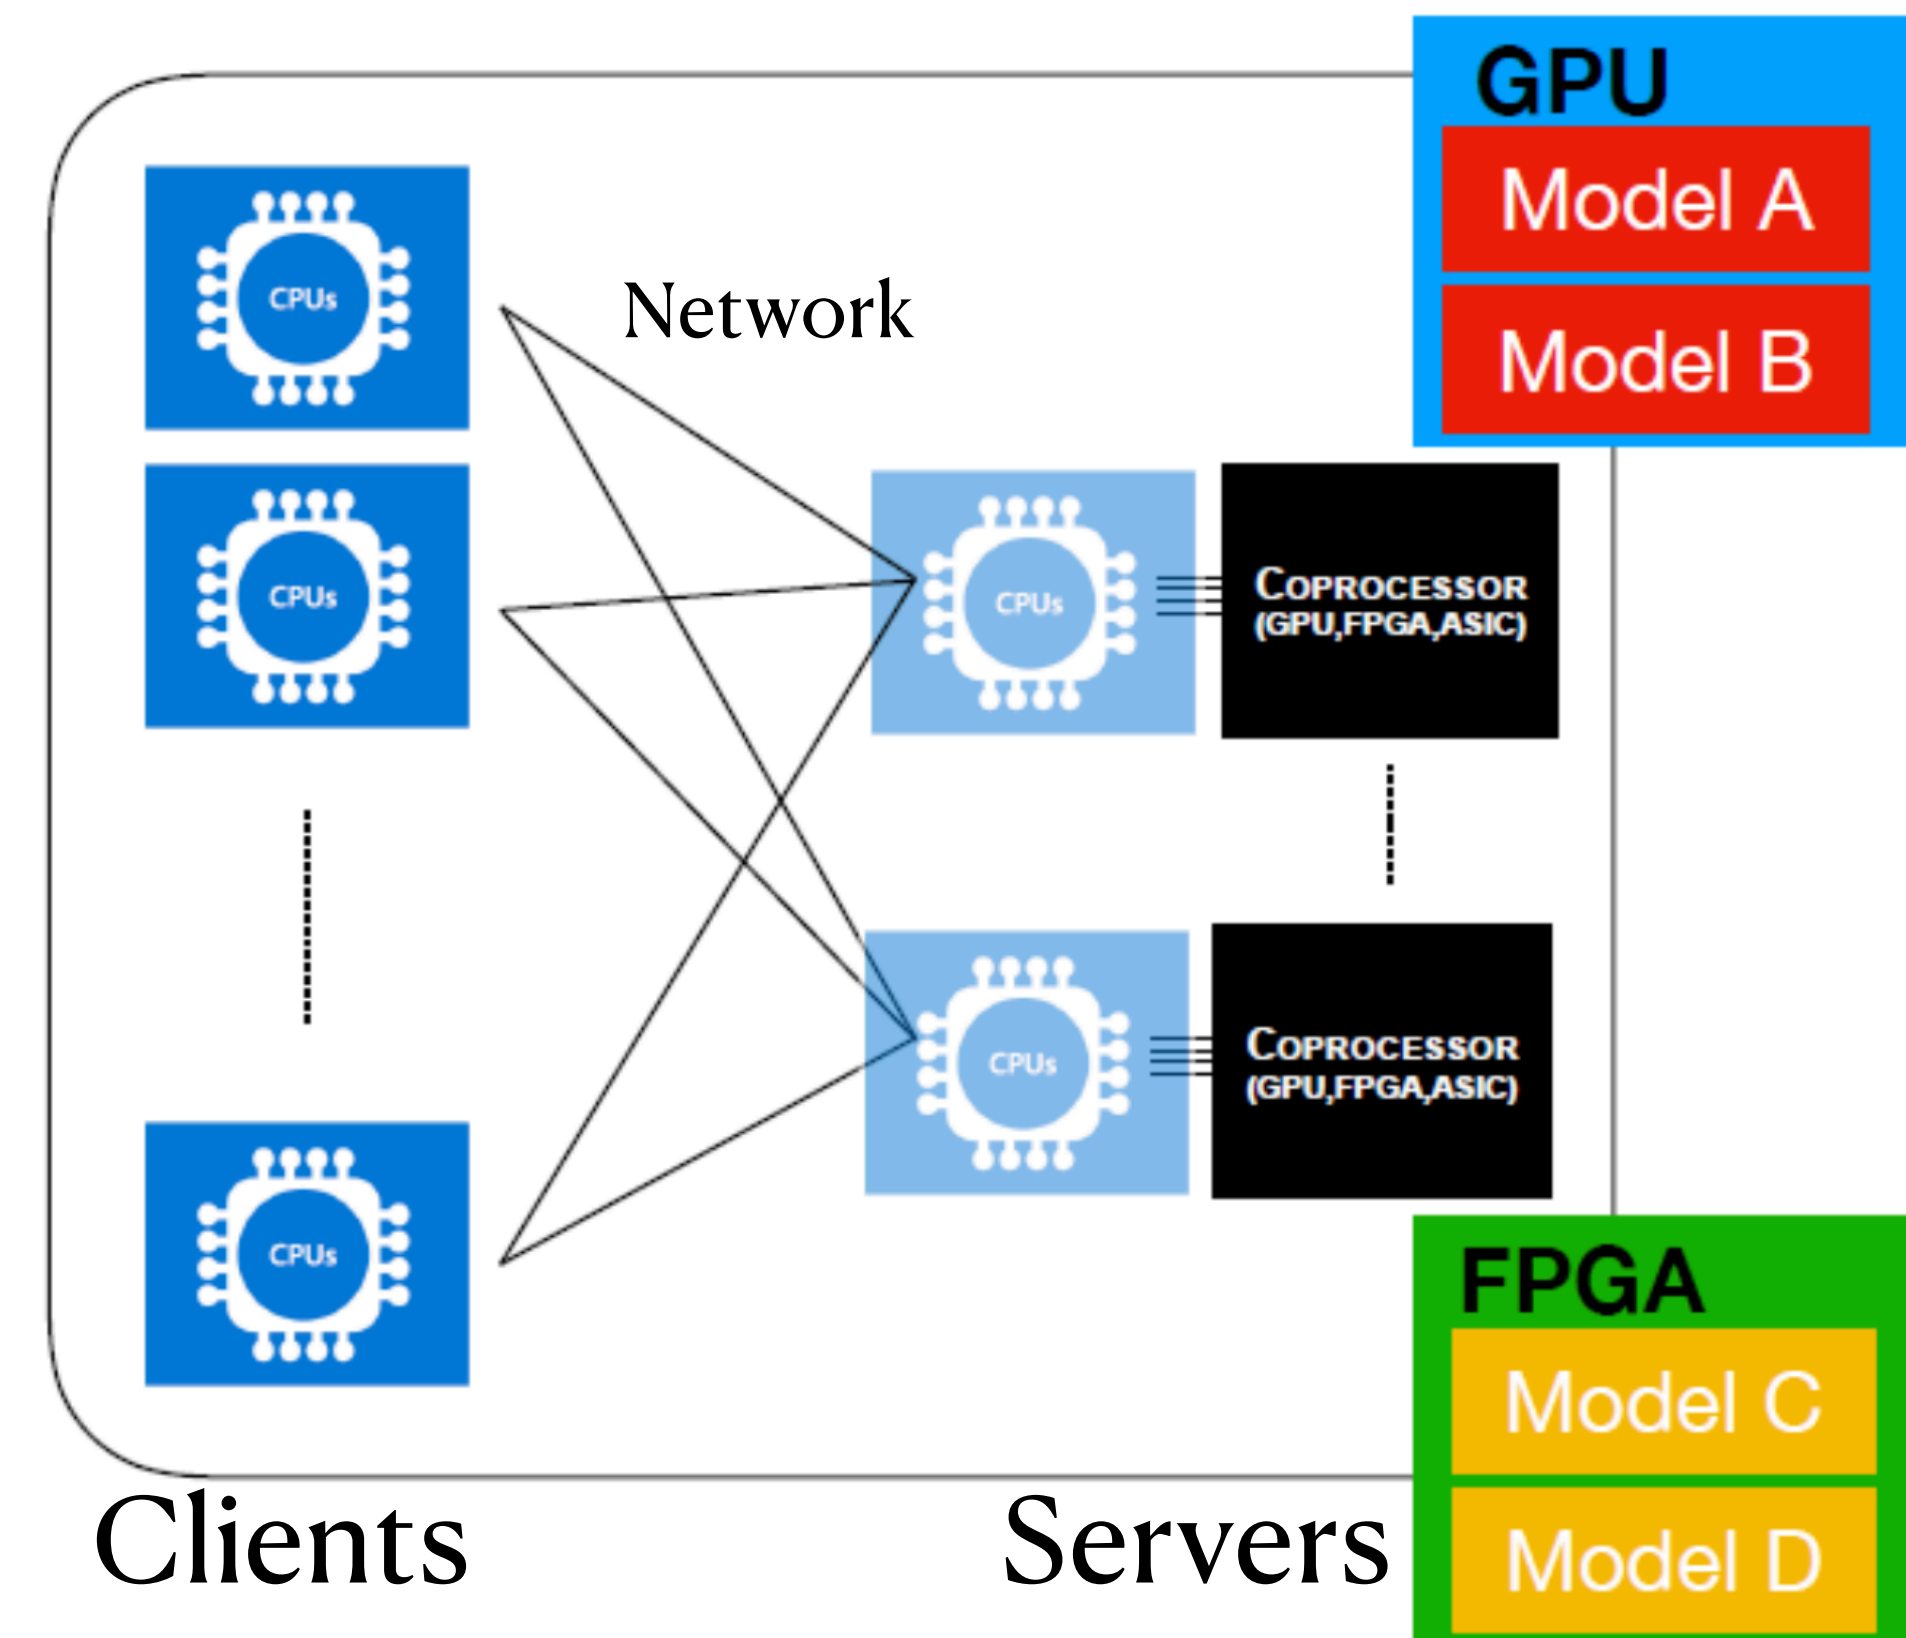
Two ML Inference Infrastructures:

- **Directly connect CPUs and coprocessors**
  - ❖ Inference running on the coprocessors directly connected to the CPU
  - ❖ Simple connection; no network load
- **Inference as a service (aaS)**
  - ❖ Clients communicate with the server, prepare the model inputs to the server and receive model outputs from the server
  - ❖ Server directs the coprocessor for model inference

## Direct



## as a Service



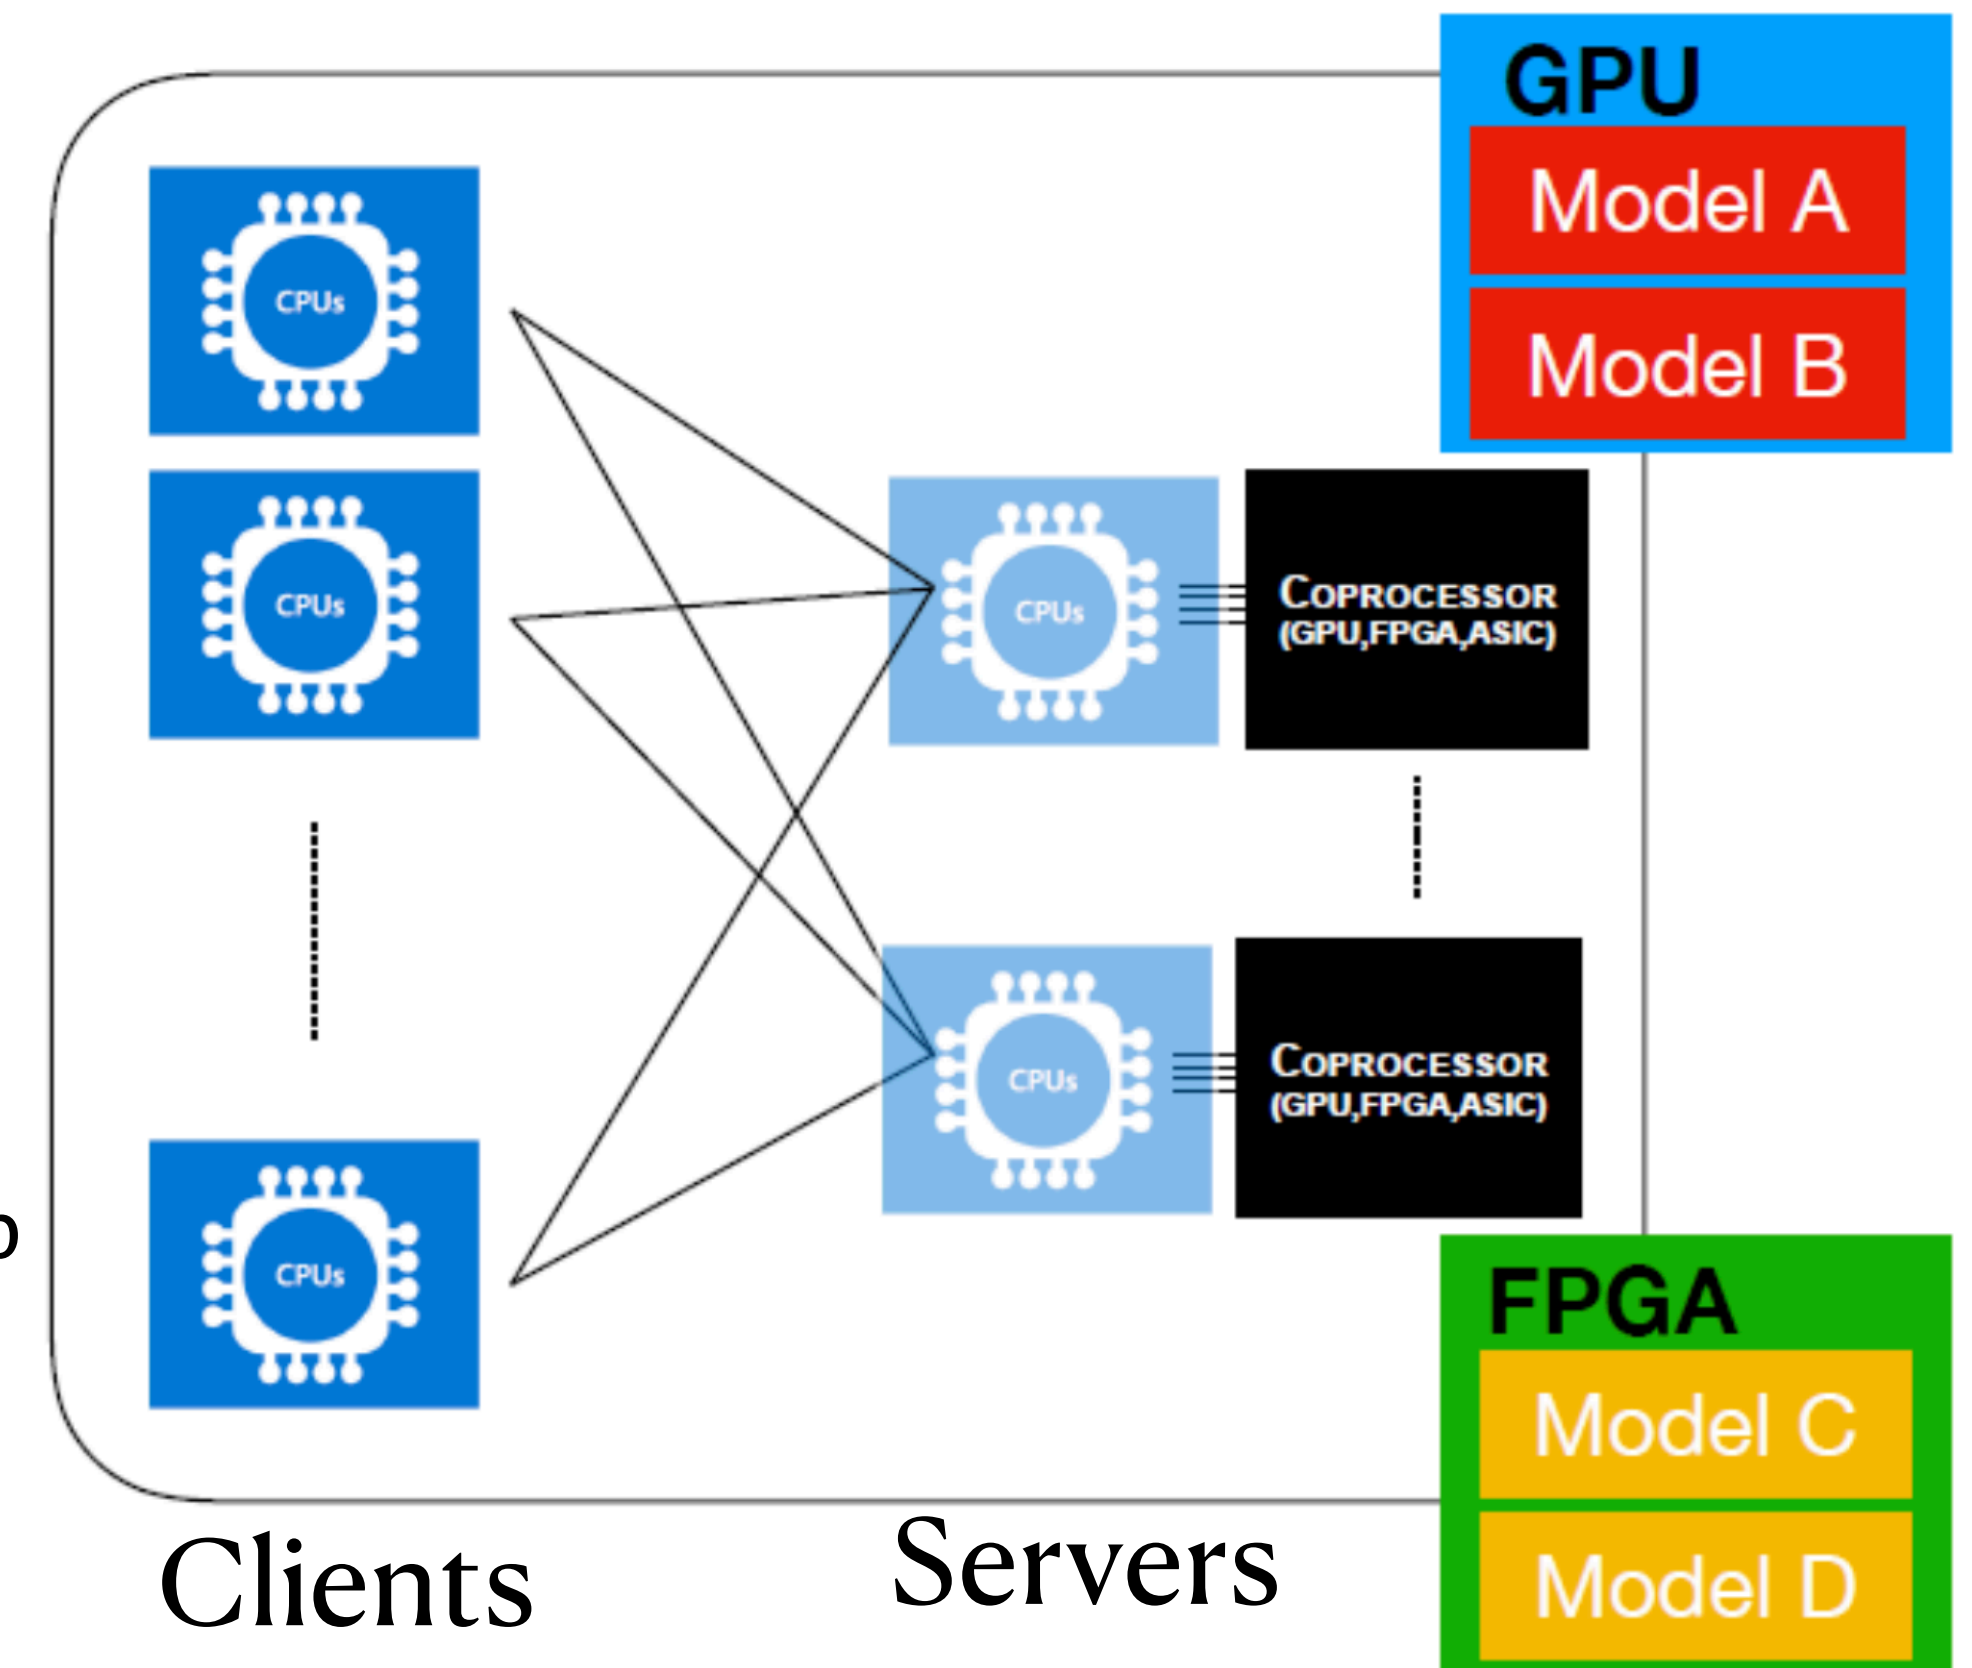


# ML Inference aaS

Benefits of running Inference aaS:

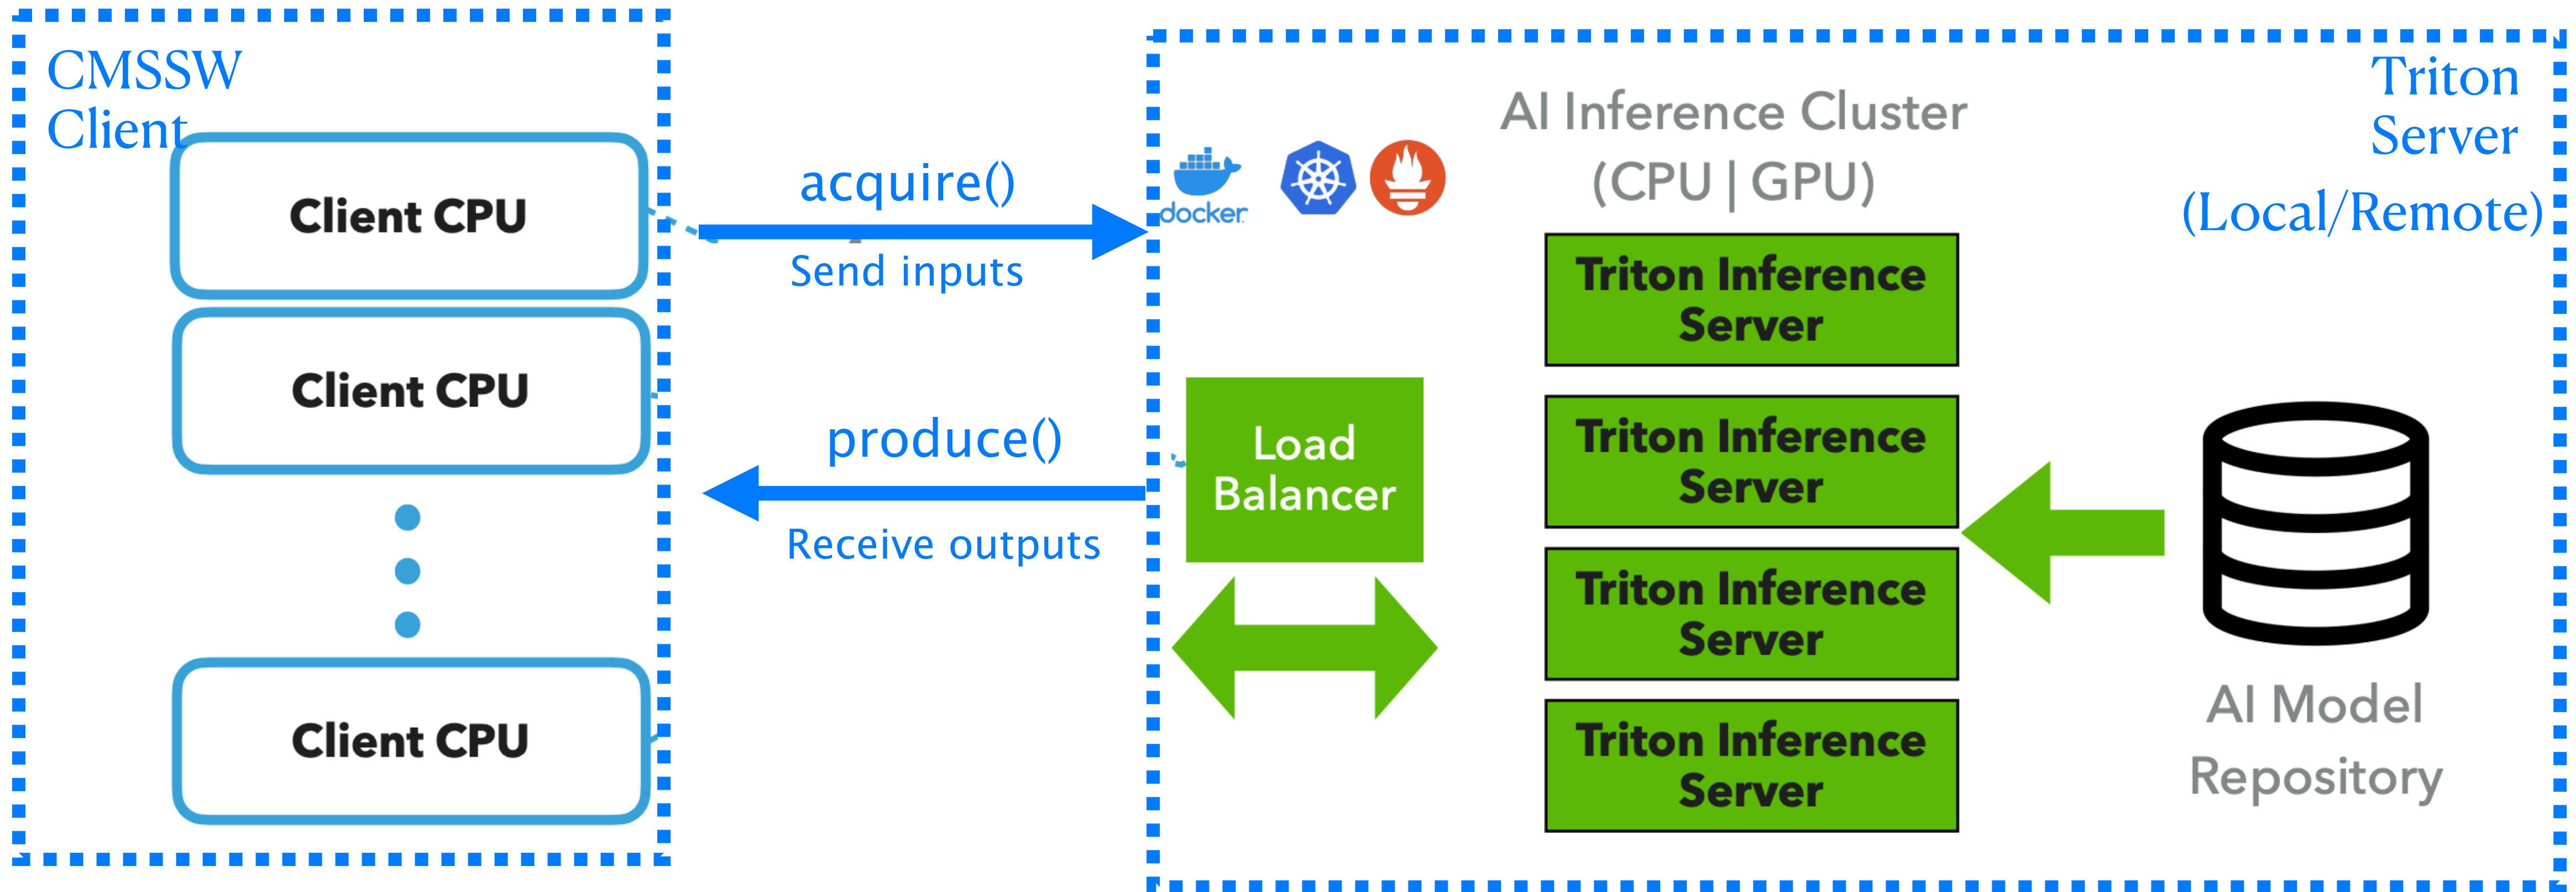
- **Factorize the ML framework out of CMSSW**
  - ✦ Only need to handle input and output conversions on the client side (i.e., in CMSSW). Different frameworks supported on the server side.
- **Simple support for different coprocessors:**
  - ✦ No need to rewrite algorithms in coprocessor-specific languages
- **More flexibility**
  - ✦ One coprocessor can serve many CPU clients
  - ✦ ML models can be deployed on different coprocessors simultaneously; choose the best coprocessor for each specific job
- **Can be more efficient in the direct connect case, with GPU shared memory**
- **Is the only option to access remote GPUs (Currently no other way to run GPU code if no local GPU is available)**

**as a Service**



Existing tools available from industry, cloud

# SONIC in CMSSW



- SONIC (Service for Optimized Network Inference on Coprocessors) available in CMSSW
- The Client in CMSSW sends the inference request with inputs for the model, and receives the outputs from the server
- NVIDIA Triton server runs the inference

# SONIC Framework in CMSSW

- **SonicCore ([repo](#))**

- ✧ Modules (EDProducer, EDFilter, EDAnalyzer) and client based classes
- ✧ Synchronous and Asynchronous modes for clients

- **SonicTriton ([repo](#))**

- ✧ Modules, clients, data types, and services for Triton inference server
- ✧ [cmsTriton](#) script to launch and manage the Triton server via Docker or Singularity

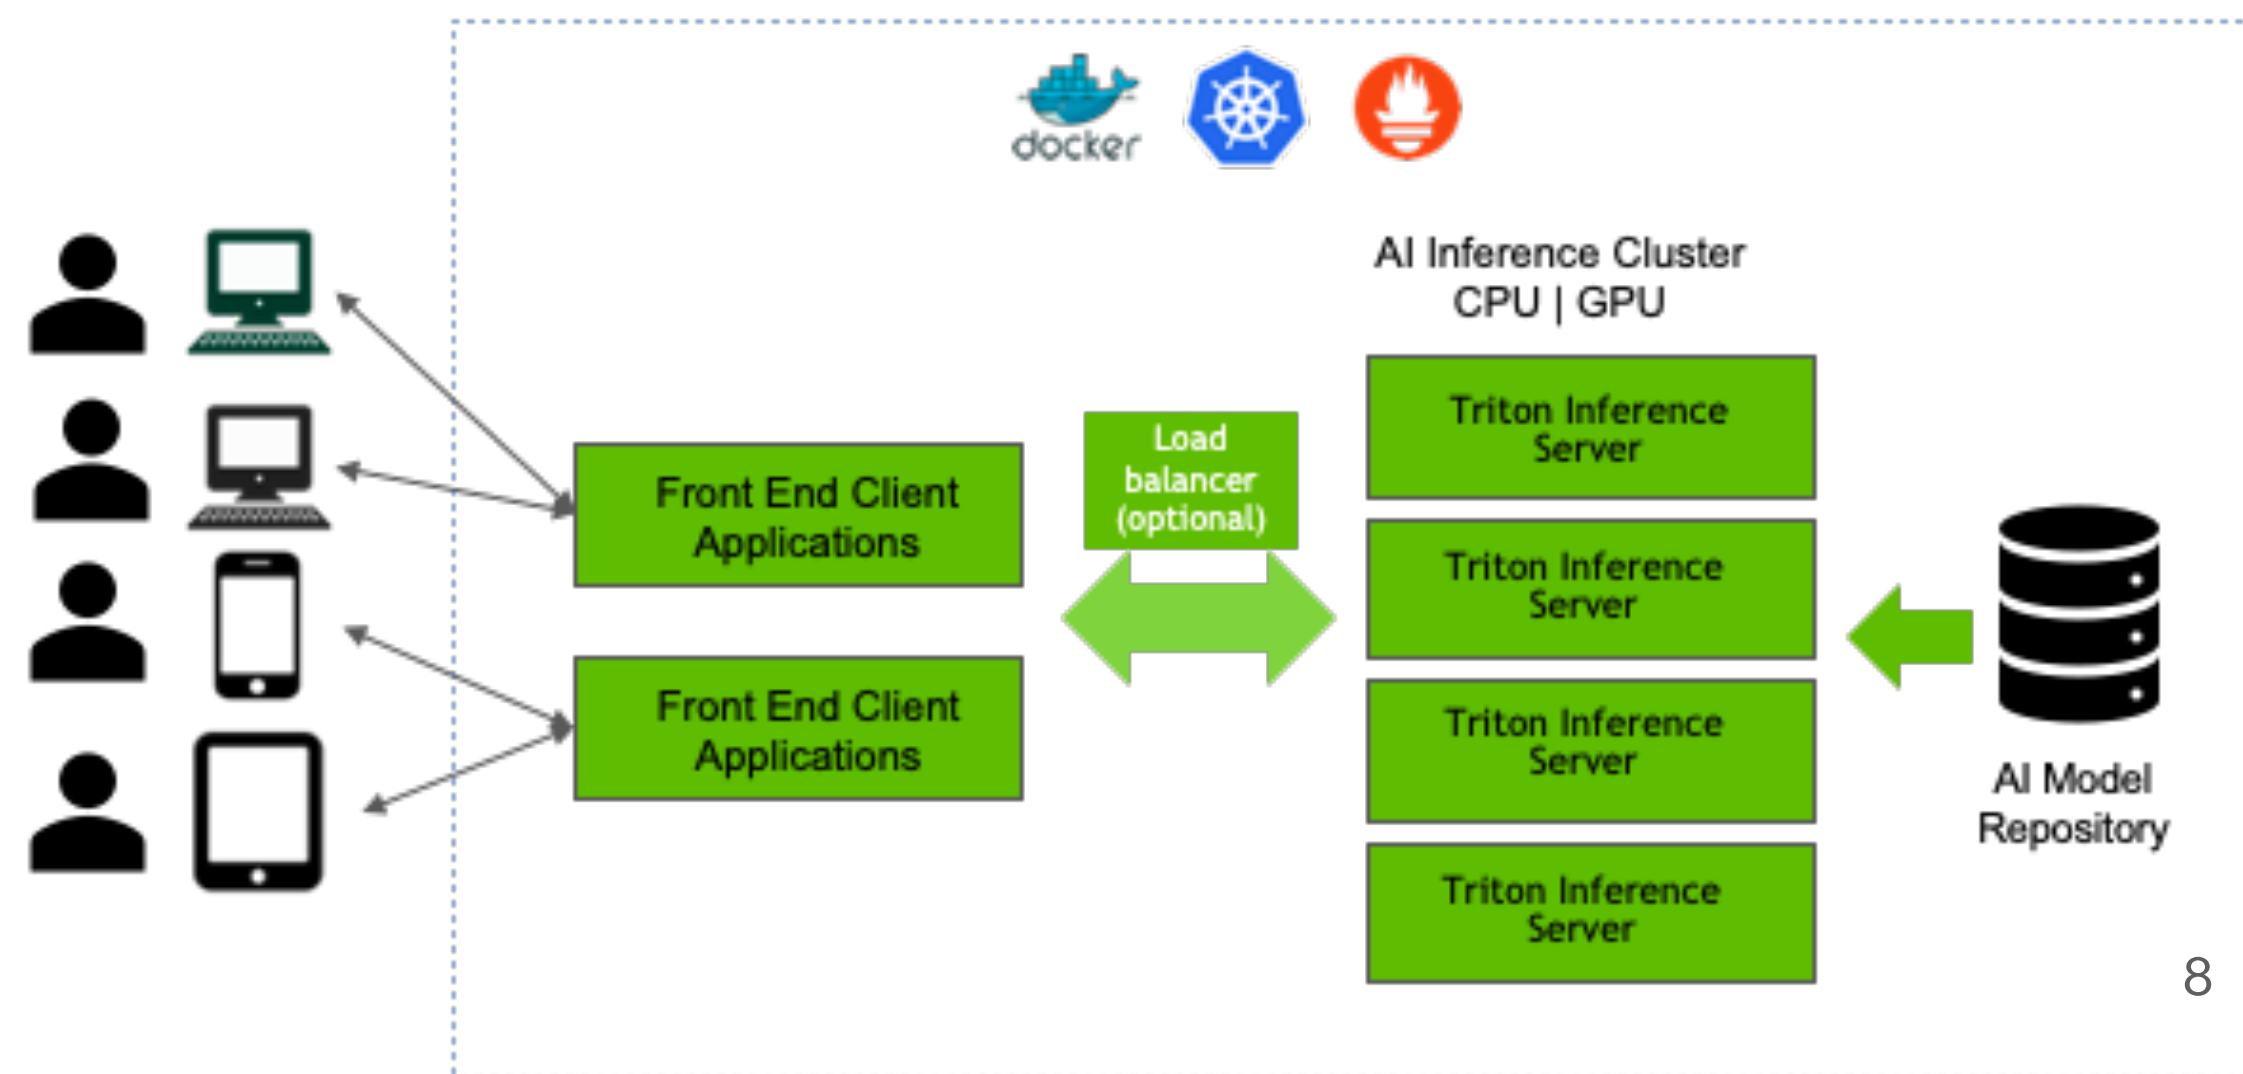
- Examples: ResNet50 and Graph Attention Network available ([code](#))

- Besides the Triton server with CPU/GPU, we are also developing the FPGA-as-a-Service Toolkit (FaaSST) for FPGAs, which SONIC also supports.

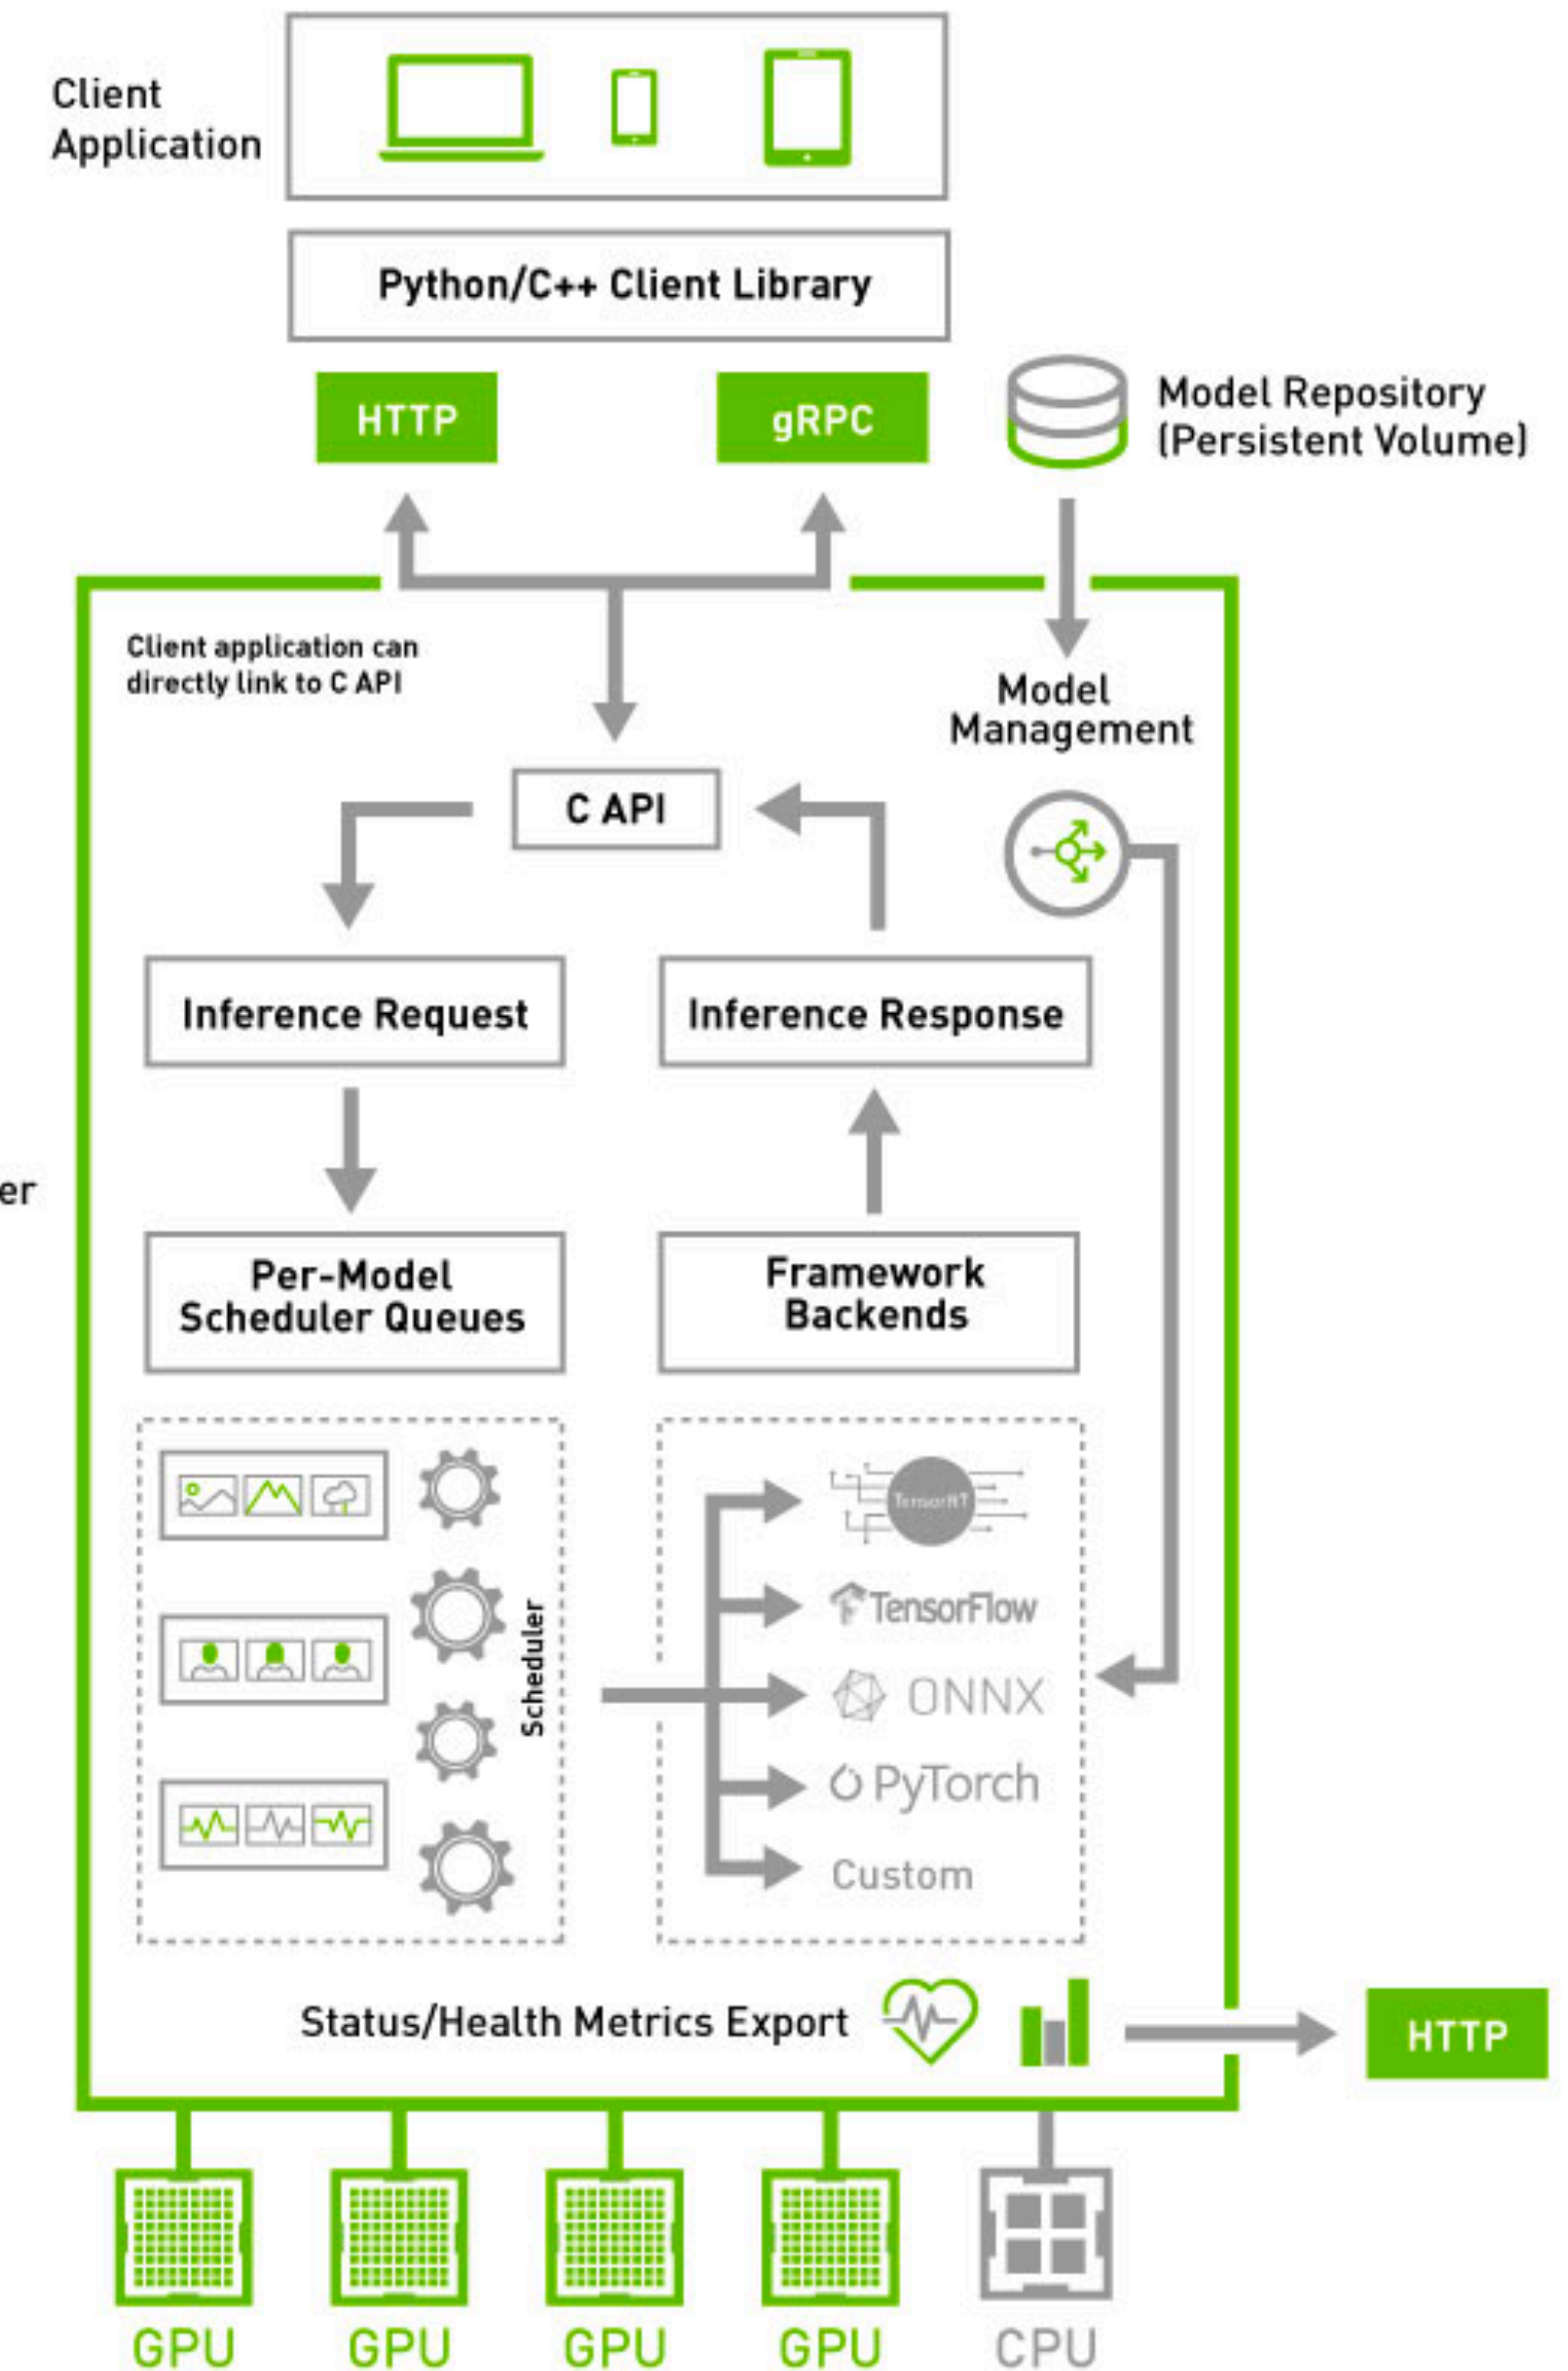


# NVIDIA Triton Inference Server

- Open source inference serving software ([webpage](#), [repo](#)) that supports numerous backends:
  - ❖ TensorFlow, Pytorch, TensorRT, ONNX for ML
  - ❖ custom backends for e.g., non-ML algorithms (cpp, CUDA, python)
- Attractive features:
  - ❖ Dynamic batching: accumulate requests from multiple events, and then process together
  - ❖ Multiple-GPU load balancing

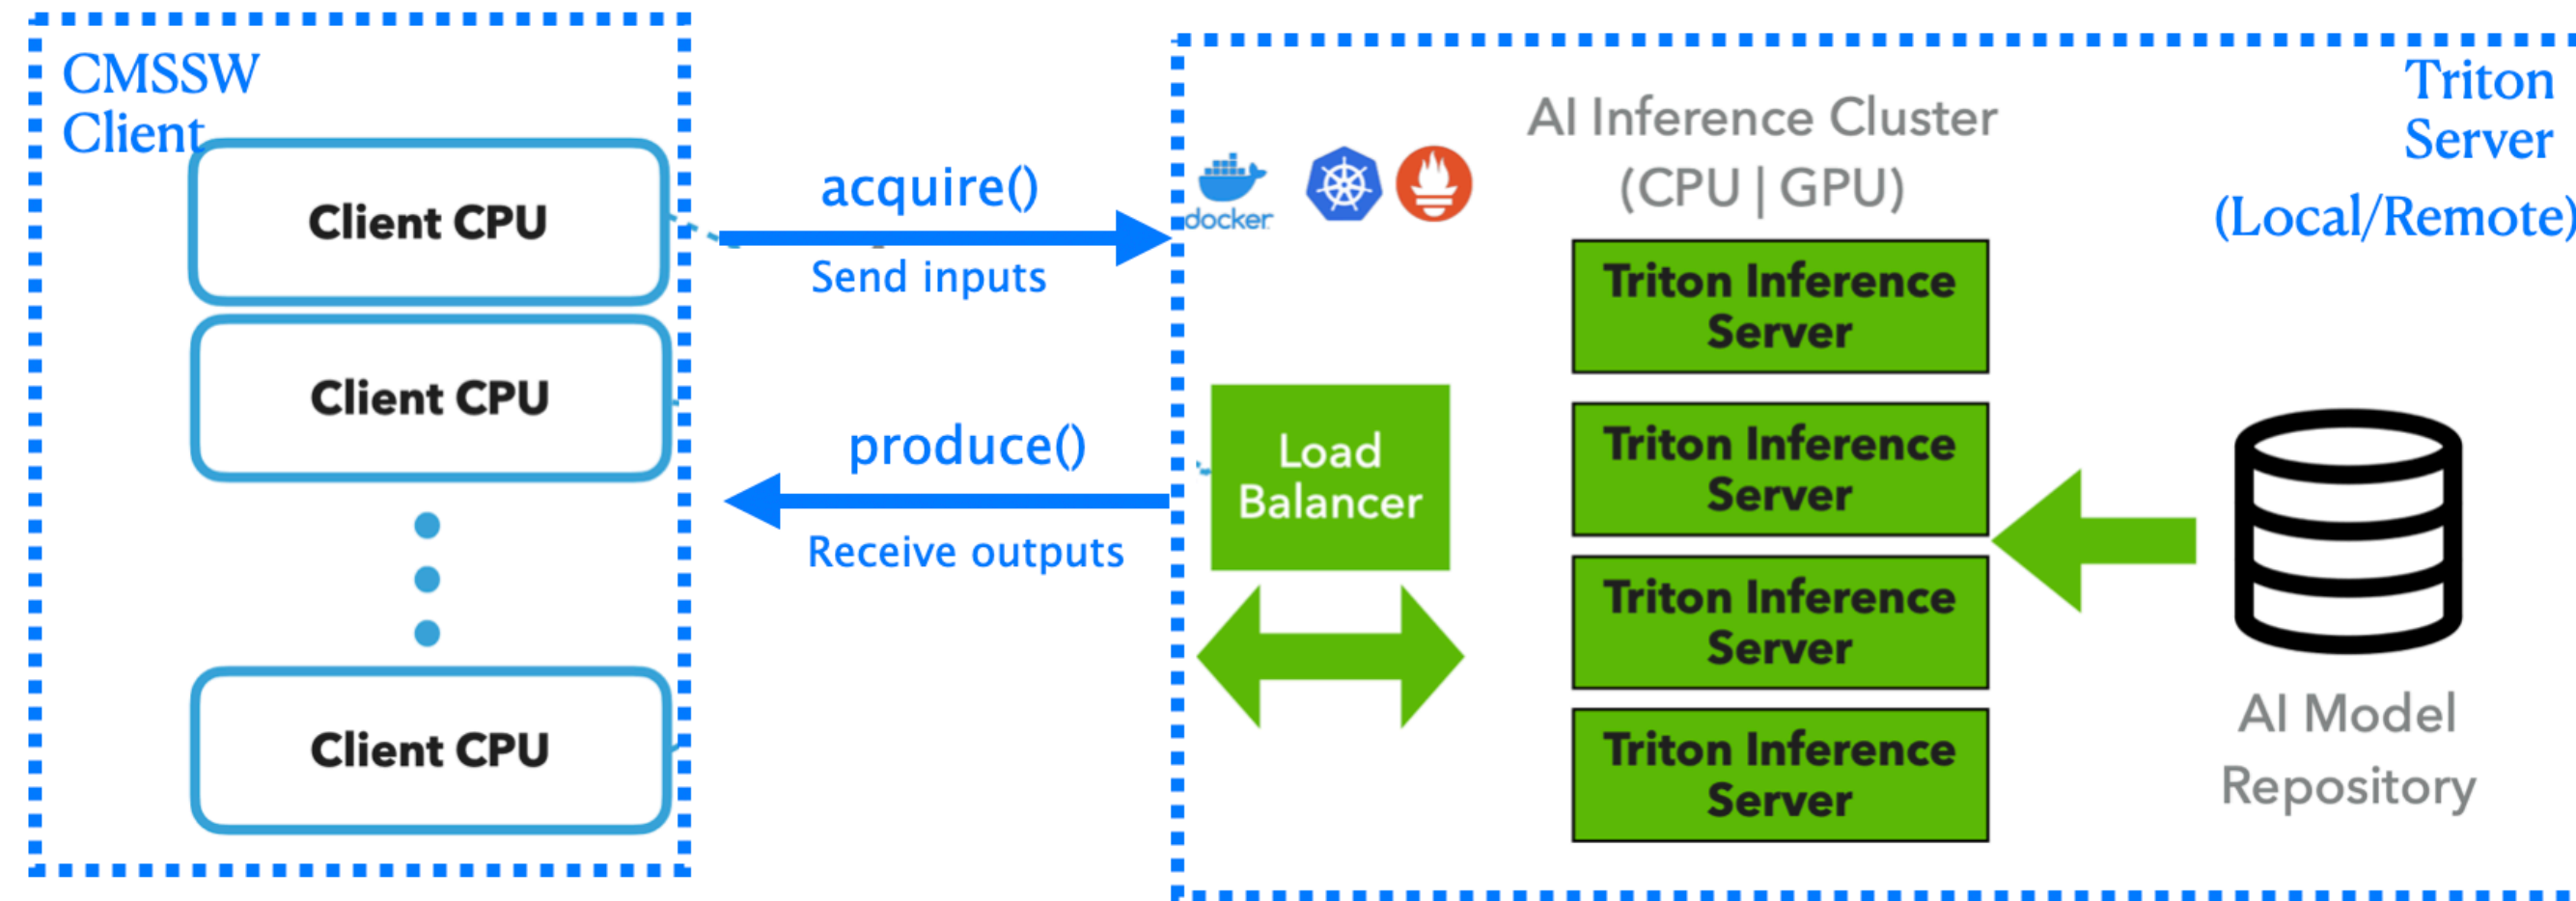


NVIDIA Triton Inference Server





# Useful Features



Useful features:

- In SONIC Triton producer, the `acquire()` function prepare inputs to the triton server, and `produce()` obtains output data from the server
- Input and output tensor shapes can be variable
- TritonService: track all available servers and models
- Local fallback server: automatically launch the local fallback server if the model is not available in the remote servers; can use either CPU or GPU (if available)
- ProcessModifier `enableSonicTriton` to turn on these features

# Useful Links

- Document on the NVIDIA Triton inference server:
  - ❖ [https://docs.nvidia.com/deeplearning/triton-inference-server/archives/triton\\_inference\\_server\\_230/user-guide/docs/](https://docs.nvidia.com/deeplearning/triton-inference-server/archives/triton_inference_server_230/user-guide/docs/)
- SONIC Core:
  - ❖ <https://github.com/cms-sw/cmssw/tree/master/HeterogeneousCore/SonicCore>
- SONIC Triton:
  - ❖ <https://github.com/cms-sw/cmssw/tree/master/HeterogeneousCore/SonicTriton>
- cmsTrion script to launch the triton server:
  - ❖ <https://github.com/cms-sw/cmssw/blob/master/HeterogeneousCore/SonicTriton/scripts/cmsTrion>
- SONIC + Triton examples:
  - ❖ <https://github.com/cms-sw/cmssw/tree/master/HeterogeneousCore/SonicTriton/test>